

1. Overview

This notebook focuses on preprocessing tweet data for sentiment analysis and emotion classification. The primary tasks include text cleaning, tokenization, and encoding categorical features. The dataset undergoes transformations using TF-IDF vectorization and OneHotEncoding to prepare it for machine learning models like Logistic Regression, Random Forest, and XGBoost.

The notebook also handles missing values, standardizes text by removing stopwords, and lemmatizes the tokens. Finally, the processed datasets are saved for future modeling and analysis.

2. Data Preparation

2.1 Importing Necessary Libraries

```
In [1]: 1 from collections import Counter
2 from nltk.stem import WordNetLemmatizer
3 from sklearn.preprocessing import LabelEncoder
4 from nltk.corpus import stopwords
5 from nltk.tokenize import word_tokenize
6 from sklearn.model_selection import train_test_split, GridSearchCV
7 from sklearn.preprocessing import LabelEncoder, OneHotEncoder, FunctionTransformer
8 from sklearn.compose import ColumnTransformer
9 from sklearn.feature_extraction.text import TfidfVectorizer
10 from sklearn.pipeline import Pipeline
11 from sklearn.metrics import classification_report, confusion_matrix
12 from sklearn.linear_model import LogisticRegression
13 from sklearn.ensemble import RandomForestClassifier
14 from xgboost import XGBClassifier
15 from tensorflow.keras.models import Sequential
16 from tensorflow.keras.layers import Embedding, LSTM, Bidirectional, Dense,
17 from tensorflow.keras.utils import to_categorical
18 from sklearn.preprocessing import OneHotEncoder
19 from nltk.corpus import wordnet
20 import category_encoders as ce
21 import pickle
22 import pandas as pd
23 import numpy as np
24 import re
25 import seaborn as sns
26 import nltk
27 import matplotlib.pyplot as plt
28 import warnings
29
30 nltk.download('stopwords')
31 nltk.download('punkt')
32 nltk.download('wordnet')
33
34 # Suppress all warnings
35 warnings.simplefilter('ignore')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Usuario\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to C:\Users\Usuario\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\Usuario\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
In [2]: 1 pd.set_option('display.max_colwidth', 1000)
```

2.2 Functions

In [3]:

```
1 def generalize_tweets(tweet_text):
2     """
3     Identify if the tweet is about a Google or Apple product, and replace
4     with 'tecproduct'.
5
6     Parameters:
7     tweet_text (str): The text of the tweet.
8
9     Returns:
10    str: 'Google' if the tweet mentions a Google product, 'Apple' if the t
11        'Both' if the tweet mentions both, 'Unknown' if it mentions neith
12    """
13    google_keywords = ['google', 'pixel', 'pixels', 'nexus', 'nexuses', 'a
14                       'chromebook', 'chromebooks', 'nest', 'nests', 'stad
15    apple_keywords = ['apple', 'apples', 'iphone', 'iphones', 'ipad', 'ipa
16                     'macbooks', 'imac', 'imacs', 'watch', 'watches', 'ai
17                     'appstore', 'ios', 'itunes']
18
19    # Ensure tweet_text is a string
20    if not isinstance(tweet_text, str):
21        return 'Unknown'
22
23    # Replace "app store" with "appstore" before tokenization
24    tweet_text = tweet_text.replace("app store", "appstore")
25
26    # Replace any occurrences of google_keywords and apple_keywords with '
27    for keyword in google_keywords + apple_keywords:
28        tweet_text = re.sub(rf'\b{keyword}\b', 'tecproduct', tweet_text, f
29
30    # Replace @ followed by any text or numbers with 'user'
31    tweet_text = re.sub(r'@\w+', 'user', tweet_text)
32
33    # Remove # in front of tecproduct if there is
34    tweet_text = re.sub(r'#tecproduct', 'tecproduct', tweet_text)
35
36    # Replace # followed by any text or numbers with 'trend'
37    tweet_text = re.sub(r'#\w+', 'trend', tweet_text)
38
39    # Remove URLs
40    tweet_text = re.sub(r'http\S+|www\S+|https\S+', 'urls', tweet_text, fl
41
42    # Rename 1g, 2g, 3g, 4g, 5g, 6g, to 'monetwork'
43    tweet_text = re.sub(r'\dg', 'monetwork', tweet_text)
44
45    return tweet_text
46
```

```
In [4]: 1 def preprocess_product_mention(df):
2         """
3         Transforming product_mention to string
4         """
5         df['product_mention'] = df['product_mention'].astype(str)
6         return df
```

3. Code

3.1 Import the database

```
In [5]: 1 # Load and prepare data
2         df = pd.read_csv('..\data\df_tweets_clean.csv')
```

3.2 Highlighting tech products

Let's make the tweets lowercase

```
In [6]: 1 df['tweet_text'] = df['tweet_text'].str.lower()
2
3         df.head()
```

Out[6]:

	tweet_text	emotion_type	product_mention
0	.@wesley83 i have a 3g iphone. after 3 hrs tweeting at #rise_austin, it was dead! i need to upgrade. plugin stations at #sxsw.	Not Positive emotion	Apple
1	@jessedee know about @fludapp ? awesome ipad/iphone app that you'll likely appreciate for its design. also, they're giving free ts at #sxsw	Positive emotion	Apple
2	@swonderlin can not wait for #ipad 2 also. they should sale them down at #sxsw.	Positive emotion	Apple
3	@sxsw i hope this year's festival isn't as crashy as this year's iphone app. #sxsw	Not Positive emotion	Apple
4	@sxtxstate great stuff on fri #sxsw: marissa mayer (google), tim o'reilly (tech books/conferences) & matt mullenweg (wordpress)	Positive emotion	Google

We have decided to rename all the technical products in the tweet_text column to a general name called tecproduct. As well as substituting all tags in a tweet with a generic name called User, and the # for another generic name called trend. Lastly, we have also replace all urls with a generic name called url. In this way, we will be able to have a more generalized tweet.

```
In [7]: 1 df['tweet_text'] = df['tweet_text'].map(generalize_tweets)
```

```
In [8]: 1 df.head()
```

Out[8]:

	tweet_text	emotion_type	product_mention
0	.user i have a monetwork tecproduct. after 3 hrs tweeting at trend, it was dead! i need to upgrade. plugin stations at trend.	Not Positive emotion	Apple
1	user know about user ? awesome tecproduct/tecproduct app that you'll likely appreciate for its design. also, they're giving free ts at trend	Positive emotion	Apple
2	user can not wait for tecproduct 2 also. they should sale them down at trend.	Positive emotion	Apple
3	user i hope this year's festival isn't as crashy as this year's tecproduct app. trend	Not Positive emotion	Apple
4	user great stuff on fri trend: marissa mayer (tecproduct), tim o'reilly (tech books/conferences) & amp; matt mullenweg (wordpress)	Positive emotion	Google

3.3 Text cleaning

3.3.1 Stop Words

Let's now proceed to eliminate the stopwords

```
In [9]: 1 stopwords_to_remove = stopwords.words('english')
2
3 df['tweet_text'] = df['tweet_text'].map(lambda x: ' '.join([word for word
4
5 df.head()
```

Out[9]:

	tweet_text	emotion_type	product_mention
0	.user monetwork tecproduct. 3 hrs tweeting trend, dead! need upgrade. plugin stations trend.	Not Positive emotion	Apple
1	user know user ? awesome tecproduct/tecproduct app likely appreciate design. also, they're giving free ts trend	Positive emotion	Apple
2	user wait tecproduct 2 also. sale trend.	Positive emotion	Apple
3	user hope year's festival crashy year's tecproduct app. trend	Not Positive emotion	Apple
4	user great stuff fri trend: marissa mayer (tecproduct), tim o'reilly (tech books/conferences) & amp; matt mullenweg (wordpress)	Positive emotion	Google

We will now proceed to remove strange characters and punctuation

```
In [10]: 1 # Remove strange characters and punctuation
2 strange_chars = '!"$%&\'()*+,-./:;<=>?[\\]^_`{|}~“!#•0ªâçûïðóôêâîô%%±°¤|~ä
3
4 df['tweet_text'] = df['tweet_text'].map(lambda x: x.translate(str.maketrans
5
6 df.head()
```

Out[10]:

	tweet_text	emotion_type	product_mention
0	user monetwork tecproduct 3 hrs tweeting trend dead need upgrade plugin stations trend	Not Positive emotion	Apple
1	user know user awesome tecproduct tecproduct app likely appreciate design also they re giving free ts trend	Positive emotion	Apple
2	user wait tecproduct 2 also sale trend	Positive emotion	Apple
3	user hope year s festival crashy year s tecproduct app trend	Not Positive emotion	Apple
4	user great stuff fri trend marissa mayer tecproduct tim o reilly tech books conferences amp matt mullenweg wordpress	Positive emotion	Google

Now we shall proceed to eliminate numbers

```
In [11]: 1 # Remove numbers
2 df['tweet_text'] = df['tweet_text'].map(lambda x: re.sub(r'\d+', '', x))
3
4 df.head()
```

Out[11]:

	tweet_text	emotion_type	product_mention
0	user monetwork tecproduct hrs tweeting trend dead need upgrade plugin stations trend	Not Positive emotion	Apple
1	user know user awesome tecproduct tecproduct app likely appreciate design also they re giving free ts trend	Positive emotion	Apple
2	user wait tecproduct also sale trend	Positive emotion	Apple
3	user hope year s festival crashy year s tecproduct app trend	Not Positive emotion	Apple
4	user great stuff fri trend marissa mayer tecproduct tim o reilly tech books conferences amp matt mullenweg wordpress	Positive emotion	Google

We are going to eliminate letters that are on their own in each individual tweet

```
In [12]: 1 df['tweet_text'] = df['tweet_text'].map(lambda x: ' '.join([word for word
2
3 df.head()
```

Out[12]:

	tweet_text	emotion_type	product_mention
0	user monetwork tecproduct hrs tweeting trend dead need upgrade plugin stations trend	Not Positive emotion	Apple
1	user know user awesome tecproduct tecproduct app likely appreciate design also they re giving free ts trend	Positive emotion	Apple
2	user wait tecproduct also sale trend	Positive emotion	Apple
3	user hope year festival crashy year tecproduct app trend	Not Positive emotion	Apple
4	user great stuff fri trend marissa mayer tecproduct tim reilly tech books conferences amp matt mullenweg wordpress	Positive emotion	Google

3.4 Lemmatization

```
In [13]: 1 # Initialize the WordNet Lemmatizer
2 lemmatizer = WordNetLemmatizer()
3
4 # Lemmatize each word in tweet_text
5 df['tweet_text'] = df['tweet_text'].map(lambda x: ' '.join([lemmatizer.lem
6
7 df.head()
```

Out[13]:

	tweet_text	emotion_type	product_mention
0	user monetwork tecproduct hr tweeting trend dead need upgrade plugin station trend	Not Positive emotion	Apple
1	user know user awesome tecproduct tecproduct app likely appreciate design also they re giving free t trend	Positive emotion	Apple
2	user wait tecproduct also sale trend	Positive emotion	Apple
3	user hope year festival crashy year tecproduct app trend	Not Positive emotion	Apple
4	user great stuff fri trend marissa mayer tecproduct tim reilly tech book conference amp matt mullenweg wordpress	Positive emotion	Google

3.5 Tokenize

```
In [14]: 1 # Tokenize the tweet_text
2 df['tweet_text_tokenized'] = df['tweet_text'].map(lambda x: word_tokenize(x))
3
4 df.head()
```

Out[14]:

	tweet_text	emotion_type	product_mention	tweet_text_tokenized
0	user monetwork tecproduct hr tweeting trend dead need upgrade plugin station trend	Not Positive emotion	Apple	[user, monetwork, tecproduct, hr, tweeting, trend, dead, need, upgrade, plugin, station, trend]
1	user know user awesome tecproduct tecproduct app likely appreciate design also they re giving free t trend	Positive emotion	Apple	[user, know, user, awesome, tecproduct, tecproduct, app, likely, appreciate, design, also, they, re, giving, free, t, trend]
2	user wait tecproduct also sale trend	Positive emotion	Apple	[user, wait, tecproduct, also, sale, trend]
3	user hope year festival crashy year tecproduct app trend	Not Positive emotion	Apple	[user, hope, year, festival, crashy, year, tecproduct, app, trend]
4	user great stuff fri trend marissa mayer tecproduct tim reilly tech book conference amp matt mullenweg wordpress	Positive emotion	Google	[user, great, stuff, fri, trend, marissa, mayer, tecproduct, tim, reilly, tech, book, conference, amp, matt, mullenweg, wordpress]

3.6 Train test split

Let's first define our variables

```
In [15]: 1 y = df['emotion_type']
2 X = df.drop(['emotion_type'], axis=1)
```

We are going to transform the variable y to numeric. Because all models need their target to be numeric. We will use the LabelEncoder.

```
In [16]: 1 # Initialize the LabelEncoder
2 label_encoder = LabelEncoder()
3
4 # Fit the encoder and transform the 'emotion_type' column
5 df['emotion_type_encoded'] = label_encoder.fit_transform(df['emotion_type'])
6
7 # The result is a new column 'emotion_type_encoded' with numeric values
8 y = df['emotion_type_encoded']
```

```
In [17]: 1 # X is the feature set and y is the target variable
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35,
```


Let's compare the shapes of `y_test` and `y_train` to see if they're somewhat similar

```
In [18]: 1 y_test.value_counts(normalize=True)
```

```
Out[18]: 0    0.655877  
        1    0.344123  
        Name: emotion_type_encoded, dtype: float64
```

```
In [19]: 1 y_train.value_counts(normalize=True)
```

```
Out[19]: 0    0.643355  
        1    0.356645  
        Name: emotion_type_encoded, dtype: float64
```

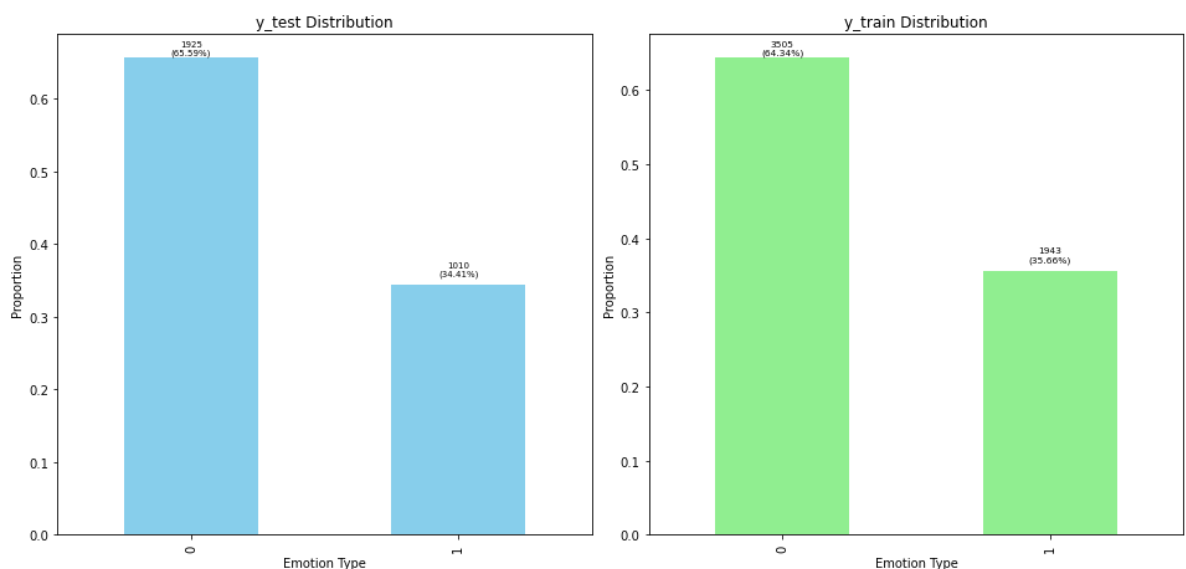
Let's see a description of the distributions

In [20]:

```

1  # Create a larger figure
2  plt.figure(figsize=(14, 7))
3
4  # y_test Distribution
5  plt.subplot(1, 2, 1)
6  y_test_counts = y_test.value_counts(normalize=True)
7  y_test_abs_counts = y_test.value_counts()
8
9  y_test_counts.plot(kind='bar', color='skyblue')
10 plt.title('y_test Distribution')
11 plt.xlabel('Emotion Type')
12 plt.ylabel('Proportion')
13
14 for i, (count, pct) in enumerate(zip(y_test_abs_counts, y_test_counts)):
15     vertical_position = pct + 0.002 if pct > 0.5 else pct + 0.01 # Small
16     plt.text(i, vertical_position, f'{count}\n({pct:.2%})', ha='center', v
17
18 # y_train Distribution
19 plt.subplot(1, 2, 2)
20 y_train_counts = y_train.value_counts(normalize=True)
21 y_train_abs_counts = y_train.value_counts()
22
23 y_train_counts.plot(kind='bar', color='lightgreen')
24 plt.title('y_train Distribution')
25 plt.xlabel('Emotion Type')
26 plt.ylabel('Proportion')
27
28 for i, (count, pct) in enumerate(zip(y_train_abs_counts, y_train_counts)):
29     vertical_position = pct + 0.002 if pct > 0.5 else pct + 0.01 # Small
30     plt.text(i, vertical_position, f'{count}\n({pct:.2%})', ha='center', v
31
32 # Adjust layout to make more space around the plots
33 plt.tight_layout()
34 plt.subplots_adjust(left=0.05, right=0.95, top=0.9, bottom=0.1)
35 plt.show()
36

```

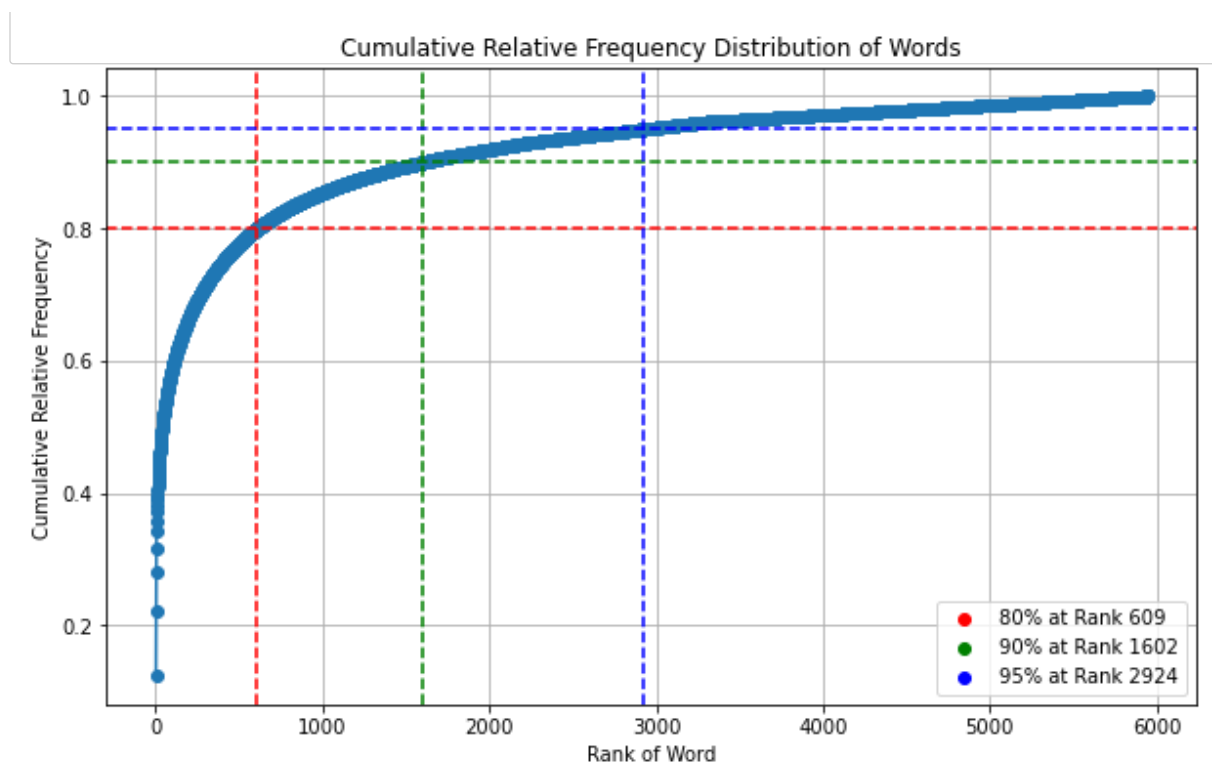


3.7 TFIDF

Before starting the procedure of vectorizing our corpus with tfidf, we want to use the a descriptive chart with the cumulative relative frequency of the words in the corpus so we can know the best threshold to set in the vectorizing process.

In [21]:

```
1 # Flatten the list of lists into a single list
2 corpus = [word for sublist in X_train['tweet_text_tokenized'] for word in
3
4 # Count the frequency of each word
5 word_freq = Counter(corpus)
6
7 # Convert the word frequency to a pandas DataFrame
8 word_freq_df = pd.DataFrame(word_freq.items(), columns=['Word', 'Frequency']
9
10 # Sort the DataFrame by frequency in descending order
11 word_freq_df = word_freq_df.sort_values(by='Frequency', ascending=False).r
12
13 # Calculate the cumulative frequency and cumulative relative frequency
14 word_freq_df['Cumulative Frequency'] = word_freq_df['Frequency'].cumsum()
15 word_freq_df['Cumulative Relative Frequency'] = word_freq_df['Cumulative F
16
17 # Find the rank where the cumulative relative frequency reaches 80%
18 rank_80_percent = word_freq_df[word_freq_df['Cumulative Relative Frequency
19 cumulative_80_percent = word_freq_df.loc[rank_80_percent - 1, 'Cumulative
20
21 # Find the rank where the cumulative relative frequency reaches 90%
22 rank_90_percent = word_freq_df[word_freq_df['Cumulative Relative Frequency
23 cumulative_90_percent = word_freq_df.loc[rank_90_percent - 1, 'Cumulative
24
25 # Find the rank where the cumulative relative frequency reaches 95%
26 rank_95_percent = word_freq_df[word_freq_df['Cumulative Relative Frequency
27 cumulative_95_percent = word_freq_df.loc[rank_95_percent - 1, 'Cumulative
28
29 # Plot the cumulative relative frequency distribution
30 plt.figure(figsize=(10, 6))
31 plt.plot(np.arange(1, len(word_freq_df) + 1), word_freq_df['Cumulative Re
32 plt.title('Cumulative Relative Frequency Distribution of Words')
33 plt.xlabel('Rank of Word')
34 plt.ylabel('Cumulative Relative Frequency')
35
36 # Mark the point where the cumulative relative frequency reaches 80%
37 plt.scatter(rank_80_percent, cumulative_80_percent, color='red', label=f'8
38 plt.axvline(x=rank_80_percent, color='red', linestyle='--')
39 plt.axhline(y=cumulative_80_percent, color='red', linestyle='--')
40
41 # Mark the point where the cumulative relative frequency reaches 90%
42 plt.scatter(rank_90_percent, cumulative_90_percent, color='green', label=f'
43 plt.axvline(x=rank_90_percent, color='green', linestyle='--')
44 plt.axhline(y=cumulative_90_percent, color='green', linestyle='--')
45
46 # Mark the point where the cumulative relative frequency reaches 95%
47 plt.scatter(rank_95_percent, cumulative_95_percent, color='blue', label=f'
48 plt.axvline(x=rank_95_percent, color='blue', linestyle='--')
49 plt.axhline(y=cumulative_95_percent, color='blue', linestyle='--')
50
51 plt.legend()
52 plt.grid(True)
53 plt.show()
54
```



As we can see, by only using 615 words, we have 80% of our entire corpus. And using 1598 words, we have a 90% of our entire corpus. With that in mind, in the past, we thought it made sense to use `max_features` of 80% in the vectorization process given that we wouldn't win much if we considered 1598 words as it only a gain of 10% more.

However, the results were not convincing and so we decided to eliminate the threshold and to get all the words. In this way we mitigated overfitting and improved the confusion matrix that is shown in notebook 03_modelling.

Knowing our `max_features` number, let's proceed and do the calculations of `tfidf`

In [22]:

```
1 # Step 1: Initialize the TfidfVectorizer with the desired max_features par
2 tfidf_vectorizer = TfidfVectorizer()
3
4 # Step 2: Fit the vectorizer on the tweet texts and transform the data int
5 tfidf_matrix_train = tfidf_vectorizer.fit_transform(X_train['tweet_text'])
6
7 # Step 3: Convert the TF-IDF matrix to a DataFrame for easier interpretati
8 # The DataFrame will have words as columns and documents as rows, with eac
9 tfidf_df_train = pd.DataFrame(tfidf_matrix_train.toarray(),
10                               columns=tfidf_vectorizer.get_feature_names_c
11                               index=X_train.index) # Keep the original ind
12
13 tfidf_df_train
```

Out[22]:

	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	abnormal	...	zimride	zing	zite
8236	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
6433	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
6987	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
952	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
5769	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
...
5734	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
5191	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
5390	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
860	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
7270	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0

5448 rows × 5918 columns

Now, let's apply the tf-idf on the X_test

```
In [23]: 1 # Step 1: Transform the tweet texts on the X_test with the vectorizer and
2         tfidf_matrix_test = tfidf_vectorizer.transform(X_test['tweet_text'])
3
4         # Step 2: Convert the TF-IDF matrix to a DataFrame for easier interpretation
5         # The DataFrame will have words as columns and documents as rows, with each
6         tfidf_df_test = pd.DataFrame(tfidf_matrix_test.toarray(),
7                                     columns=tfidf_vectorizer.get_feature_names_out(),
8                                     index=X_test.index) # Keep the original index
9
10        tfidf_df_test
```

Out[23]:

	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	abnormal	...	zimride	zing
2865	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.379029	0.0	...	0.0	0.0
8174	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
5933	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
4857	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
6858	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
...
2662	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
5958	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
5840	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
2154	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0
1344	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0

2935 rows × 5918 columns

We are first going to merge the product_mention and the y_train, y_test columns to tfidf_df_train and tfidf_df_test separately.

```
In [24]: 1 # For tfidf_df_train. Let's join tfidf_df_train and X_train by index
2 df_train = X_train[['product_mention']].join(tfidf_df_train)
3
4 # Let's now join y_train with df_train
5 df_train = df_train.join(y_train)
6
7 df_train
```

Out[24]:

	product_mention	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	...	zing	zite
8236	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
6433	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
6987	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
952	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
5769	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
5734	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
5191	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
5390	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
860	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
7270	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0

5448 rows × 5920 columns


```
In [25]: 1 # For tfidf_df_test. Let's join tfidf_df_test and X_test by index
2 df_test = X_test[['product_mention']].join(tfidf_df_test)
3
4 # Let's now join y_test with df_test
5 df_test = df_test.join(y_test)
6
7 df_test
```

Out[25]:

	product_mention	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	...	zing	z
2865	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.379029	...	0.0	(
8174	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
5933	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
4857	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
6858	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
...	
2662	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
5958	Google	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
5840	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
2154	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(
1344	Apple	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	...	0.0	(

2935 rows × 5920 columns

We shall now proceed to save the pickle of the tfidf_vectorizer

```
In [26]: 1 # Save the tfidf_vectorizer in a pickle folder called pickle_objects
2 with open('../pickle_objects/tfidf_vectorizer.pkl', 'wb') as file:
3     pickle.dump(tfidf_vectorizer, file)
```

Let's do a one hot encoder on the product_mention column of df_train

In [27]:

```
1 # Initialize the OneHotEncoder
2 onehot_encoder = OneHotEncoder(sparse=False, drop='first') # drop='first'
3
4 # Fit and transform the 'product_mention' column
5 encoded_features = onehot_encoder.fit_transform(df_train[['product_mention']])
6
7 # Convert the array back to a DataFrame
8 encoded_df = pd.DataFrame(encoded_features, columns=onehot_encoder.get_feature_names_out())
9
10 # Concatenate the encoded columns back to the original DataFrame
11 df_train_encoded = pd.concat([df_train.reset_index(drop=True), encoded_df], axis=1)
12
13 # Drop the 'product_mention' column
14 df_train_encoded.drop('product_mention', axis=1, inplace=True)
15
16 # Display the DataFrame with the encoded columns
17 df_train_encoded.head()
```

Out[27]:

	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	abnormal	...	zms	zombie	zomg
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0

5 rows × 5921 columns

Let's do the One Hot Encoder for X_test

```
In [28]: 1 # Transform the 'product_mention' column
2 encoded_features = onehot_encoder.transform(df_test[['product_mention']])
3
4 # Convert the array back to a DataFrame
5 encoded_df = pd.DataFrame(encoded_features, columns=onehot_encoder.get_fea
6
7 # Concatenate the encoded columns back to the original DataFrame
8 df_test_encoded = pd.concat([df_test.reset_index(drop=True), encoded_df],
9
10 # Drop the 'product_mention' column
11 df_test_encoded.drop('product_mention', axis=1, inplace=True)
12
13 # Display the DataFrame with the encoded columns
14 df_test_encoded.head()
```

Out[28]:

	aapl	aaron	ab	abacus	abba	abc	aber	ability	able	abnormal	...	zms	zombie	zor
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.379029	0.0	...	0.0	0.0	(
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0	(
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0	(
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0	(
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.0	0.0	(

5 rows × 5921 columns

Let's save the pickle for the One Hot Encoder

```
In [29]: 1 # Save the ohe in a pickle folder called pickle_objects
2 with open('../pickle_objects/ohe.pkl', 'wb') as file:
3     pickle.dump(onehot_encoder, file)
```

4. Export to csv

```
In [30]: 1 # Save df_train_encoded to a CSV file
2 df_train_encoded.to_csv('../data/train_processed.csv', index=False)
3 # The index=False argument ensures that the DataFrame index is not written
4
5 # Save df_test_encoded to a CSV file
6 df_test_encoded.to_csv('../data/test_processed.csv', index=False)
```